



**BITS Pilani**

Hyderabad Campus

## SS ZG 526: Distributed Computing (CS2)

### Logical Time

Prof. Geetha

Associate Professor, Dept. of Computer Sc. & Information Systems

BITS Pilani Hyderabad Campus

[geetha@hyderabad.bits-pilani.ac.in](mailto:geetha@hyderabad.bits-pilani.ac.in)

# Introduction

- ❖ **concept of causality between events** → important for design and analysis of parallel and distributed computing and operating systems
- ❖ causality is tracked using physical time
- ❖ But in distributed systems
  - ❖ not possible to have global physical time
  - ❖ possible to realize only an approximation of physical time
- ❖ way of measuring time
  - ❖ asynchronous distributed computations progress in **spurts**
  - ❖ logical time also advances in **jumps**
  - ❖ sufficient to capture the fundamental **monotonicity** property associated with causality in distributed systems

# Introduction

---

- ❖ **monotonicity property: if event  $a$  causally affects event  $b$ , then timestamp of  $a$  is smaller than timestamp of  $b$**
- ❖ Ways to represent logical time:
  - ❖ scalar time
  - ❖ vector time

# Introduction

- ❖ **Causality / Causal Precedence Relation**
- ❖ described among events in a distributed system
- ❖ useful in reasoning, analyzing, and drawing inferences about a computation
- ❖ helps to solve several problems in distributed systems
  - ❖ distributed algorithms design
  - ❖ tracking of dependent events
  - ❖ knowledge about progress
  - ❖ concurrency measure

# Introduction

## ➤ Distributed algorithms design

knowledge of the causal precedence relation among events helps to:

- ensure **liveness** and **fairness** in mutual exclusion algorithms
- maintain **consistency** in replicated databases
- design correct deadlock detection algorithms

## ➤ Tracking of dependent events

- helps construct a consistent state for resuming re-execution
- in failure recovery, helps build a **checkpoint**
- in replicated databases, aids in detection of file inconsistencies

- **Knowledge about progress - knowledge of causal dependency among events**
  - helps measure progress of processes in distributed computation
  - helps to discard obsolete information, garbage collection, and termination detection
- **Concurrency measure - knowledge regarding how many events are causally dependent**
  - useful in measuring amount of concurrency
  - events not causally related are concurrent
  - analysis of causality tells about concurrency in program

# Capturing Causality between Events

- in distributed systems
  - rate of event occurrence is of high magnitude
  - event execution time is of low magnitude
- physical clocks must be precisely synchronized
- in distributed computation
  - clocks accurate to a few tens of milliseconds are not sufficient
  - progress occurs in spurts
  - interaction between processes occurs in spurts
- logical clocks can accurately capture
  - causality relation between events produced by a program execution
  - fundamental monotonicity property

# Logical Clocks



- every process has a logical clock
- logical clock is advanced using a set of rules
- each event is assigned a timestamp
- causality relation between events can be inferred from their timestamps
- timestamps follow the monotonicity property



# Representing Logical Clocks



1. Lamport's scalar clocks
  - time is represented by non-negative integers
2. Vector clocks (Fidge, Mattern and Schmuck)
  - time is represented using a vector of non-negative integers

# Real world applications of Logical clocks



- ❖ Scalar clocks => Amazon S3
- ❖ Vector clocks => Voldemort (Linkedin) Amazon Dynamo, Version control systems etc.

Self assessment Question:

- ❖ Why Lamport's clock cannot be used in blockchain technology?

# Definition of Logical Clocks



- system of logical clocks consists of
  - time domain  $T$
  - logical clock  $C$
- elements of  $T$  form a partially ordered set over a relation  $<$
- $<$ :
  - happened before or causal precedence relation
  - analogous to 'earlier than' relation provided by physical time

# Definition of Logical Clocks



- logical clock  $C$ 
  - is a function that maps an event  $e$  in a distributed system to an element in the time domain  $T$
  - denoted as  $C(e)$
  - $C(e)$ : timestamp of  $e$
  - $C$  is defined as:  $C : H \mapsto T$
  - satisfies the following property –
    - for any 2 events  $e_i$  and  $e_j$ ,  $e_i \rightarrow e_j \implies C(e_i) < C(e_j)$ .
    - monotonicity property
    - known as *clock consistency condition*

# Definition of Logical Clocks



•system of clocks is *strongly consistent* if  $T$  and  $C$  satisfy the following:

•for 2 events  $e_i$  and  $e_j$ ,

$$e_i \rightarrow e_j \Leftrightarrow C(e_i) < C(e_j)$$

# Implementing Logical Clocks

---

- ❖ Two issues need to be addressed
  - ❖ data structures local to every process to represent logical time
  - ❖ a protocol to update the data structures to ensure the consistency condition

# Implementing Logical Clocks

## Data Structures

- Each process  $p_i$  maintains two data structures
- A *local logical clock*
  - denoted by  $lc_i$
  - helps process  $p_i$  measure its own progress
- A *logical global clock*
  - denoted by  $gc_i$
  - represents  $p_i$ 's local view of the logical global time
  - allows  $p_i$  to assign consistent timestamps to its local events
  - $lc_i$  is a part of  $gc_i$

# Implementing Logical Clocks

Protocol - ensures consistent management of:

- a process's logical clock
- process's view of global time
- consists of the following 2 rules:
  - **R1:** governs how the local logical clock is updated by a process when it executes an event (send, receive, or internal)
  - **R2:**
    - governs how a process updates its global logical clock to update its view of the global time and global progress
    - dictates what information about logical time is piggybacked in a message
    - how this information is used by the receiving process to update its view of the global time



# Implementing Logical Clocks

---

- systems of logical clocks differ in
  - representation of logical time
  - the protocol to update the logical clocks
- all logical clock systems
  - implement rules **R1** and **R2**
  - ensure the fundamental monotonicity property associated with causality
  - provide users with some additional properties

# Scalar Time - Definition

- representation was proposed by Lamport in 1978 to totally order events in distributed system
- in this representation, time domain is the set of non-negative integers
- $C_i$ :
  - integer variable
  - denotes logical local clock of  $p_i$  and its local view of global time

# Scalar Time - Definition

## Rules for updating clock:

- **R1:** Before executing an event (send, receive, or internal), process  $p_i$  executes

$$C_i := C_i + d \quad (d > 0)$$

For each execution of **R1**,

- *$d$  can have different value*
- value of  $d$  may be application-dependent

Usually  $d = 1$

- helps to identify the time of each event uniquely at a process
- keeps rate of increase of  $d$  to lowest level


# Scalar Time - Definition

## Rules for updating clock (contd. from previous slide):

**R2:** When process  $p_i$  receives a message with timestamp  $C_{msg}$ , it executes the following actions:

1.  $C_i = \max(C_i, C_{msg})$ ;
2. execute **R1**;
3. deliver the message.

# Evolution of Scalar Time



Space-time diagram of a distributed system

# Scalar Time – Basic Properties

---

- Consistency
- Total Ordering
- Event Counting
- No Strong Consistency

# Scalar Time – Basic Properties


## Consistency Property

- for two events  $e_i$  and  $e_j$ ,  $e_i \rightarrow e_j \implies C(e_i) < C(e_j)$
- scalar clocks satisfy monotonicity property
- hence, they also satisfy consistency property

# Scalar Time – Basic Properties

## Total Ordering

- Scalar clocks are used to totally order events in a distributed system
- Problem in totally ordering events – 2 or more events at different processes may have same timestamp  
for 2 events  $e_1$  and  $e_2$ ,  $C(e_1) = C(e_2) \implies e_1 \parallel e_2$





# Scalar Time – Basic Properties

## Total Ordering

- Tie breaking mechanism is needed
- Tie breaking procedure:
  - process identifiers are linearly ordered
  - break tie among events with identical scalar timestamps on the basis of their process identifiers
  - lower process identifier implies higher priority
  - timestamp of an event is denoted by a tuple  $(t, i)$ 
    - $t \rightarrow$  *time of occurrence*
    - $i \rightarrow$  *identity of the process where it occurred*

# Scalar Time – Basic Properties

## Total Ordering

- Total order relation  $<$  on 2 events  $x$  and  $y$  with timestamps  $(h, i)$  and  $(k, j)$  respectively, is defined:  
$$x < y \Leftrightarrow h < k \text{ or } (h = k \text{ and } i < j)$$
- events that occur at the same logical scalar time are independent (i.e., they are not causally related)
- such events can be ordered using any arbitrary criterion without violating the causality relation  $\rightarrow$
- so, total order is consistent with the causality relation “ $\rightarrow$ ”
- Note:-  $x < y \Rightarrow x \rightarrow y \vee x || y$

# Scalar Time – Basic Properties

---

## Total Ordering


- Utility
  - total order is generally used to ensure liveness properties in distributed algorithms
    - Liveness property:
      - something good eventually happens
      - system makes progress, no starvation, programs terminate
  - requests are timestamped and served according to the total order based on these timestamps

# Scalar Time – Basic Properties

## Event Counting

- If the increment value  $d = 1$ , scalar time has the following property:
  - if event  $e$  has a timestamp  $h$ , then  $h - 1$ 
    - is the minimum logical duration required before producing event  $e$
    - counted in units of events
    - called the ***height*** of event  $e$
  - implies that  $h - 1$  events have been produced sequentially before event  $e$ , irrespective of the processes that produced these events

# Scalar Time – Basic Properties




- height of event  $a = 4 - 1 = 3$ , i.e., 3 events precede  $a$  on the longest causal path ending at  $a$
- height of event  $b = 6 - 1 = 5$ , i.e., 5 events precede  $b$  on the longest causal path ending at  $b$

# Scalar Time – Basic Properties

## No strong consistency

- system of scalar clocks is not strongly consistent
- for 2 events  $e_i$  and  $e_j$ ,  $C(e_i) < C(e_j) \not\Rightarrow e_i \rightarrow e_j$




- scalar timestamp of 3rd event of  $p_1 <$  scalar timestamp of 3rd event of  $p_2$
- but the former did not happen before the latter

# Scalar Time – Basic Properties

## No strong consistency

- scalar clocks are not strongly consistent because logical local clock and logical global clock of a process are combined into one
- results in the loss causal dependency information among events at different processes



- when  $p_2$  receives the 1st message from  $p_1$ ,  $p_2$  updates its clock to 3
- forgets that the timestamp of the latest event at  $p_1$  on which it depends is 2

# Vector Time - Definition

- time domain is represented by a set of  $n$ -dimensional non-negative integer vectors
- each process  $p_i$  maintains a vector  $vt_i[1..n]$ 
  - $vt_i[i]$  : *local logical clock of  $p_i$*
  - describes progress of logical time at  $p_i$
  - $vt_i[j]$ : *represents process  $p_i$ 's latest knowledge of process  $p_j$ 's local time*
  - if  $vt_i[j] = x$ 
    - $p_i$  *knows that local time at  $p_j$  has progressed till  $x$*
- vector  $vt_i$  constitutes  $p_i$ 's view of the global logical time
- $vt_i$  *is used to timestamp events*



# Vector Time - Definition

Rules used by  $p_i$  for clock updating

- **R1:** Before executing an event, process  $p_i$  updates its local logical time as follows:

$$vt_i[i] = vt_i[i] + d, \quad d > 0$$

# Vector Time - Definition

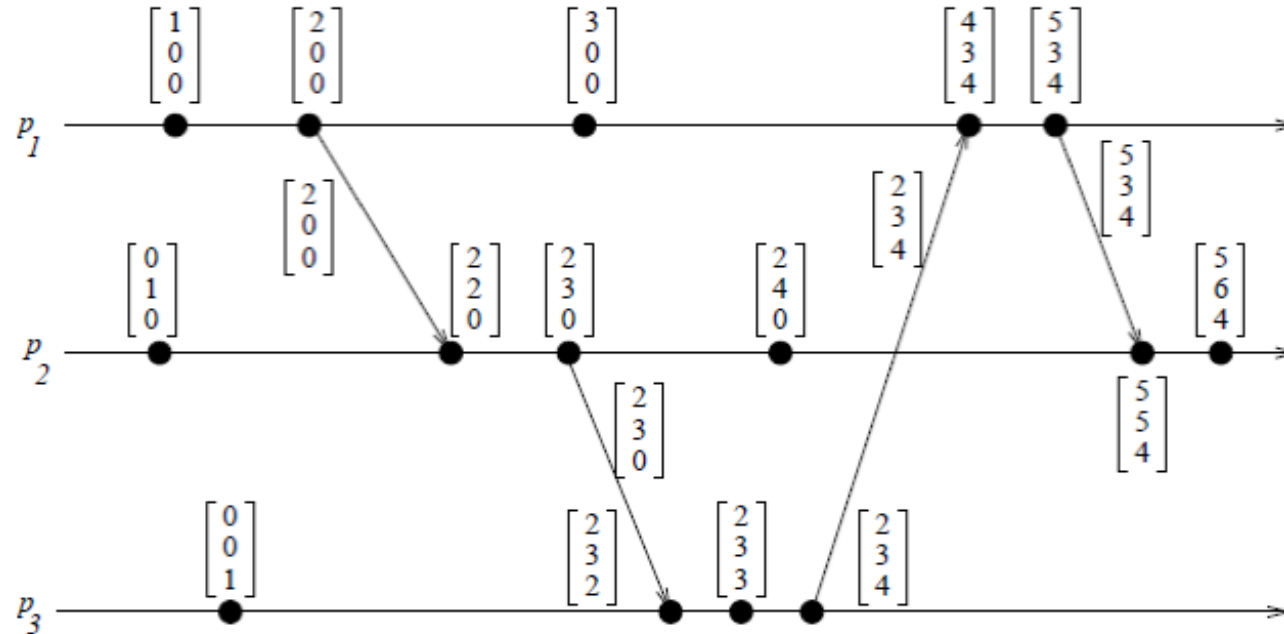
## Rules used by $p_i$ for clock updating

- **R2:**
  - each message  $m$  is piggybacked with the vector clock  $vt$  of the sender process at sending time
  - when  $p_i$  receives message  $(m, vt)$ , it executes the following sequence of actions:
    1. update its global logical time as:
$$1 \leq k \leq n: vt_i[k] = \max(vt_i[k], vt[k]);$$
    2. execute R1;
    3. deliver the message  $m$ .

# Vector Time - Definition

## Rules used by $p_i$ for clock updating

- **R2:**
  - timestamp associated with an event is the value of the vector clock of its process when the event is executed
  - initially, a vector clock is  $[0, 0, 0, \dots, 0]$



vector clock's  
progress with  $d = 1$

# Vector Time - Definition

Relations for comparing 2 vector timestamps  $vh$  and  $vk$

- $vh = vk \Leftrightarrow \forall x : vh[x] = vk[x]$
- $vh \leq vk \Leftrightarrow \forall x : vh[x] \leq vk[x]$
- $vh < vk \Leftrightarrow vh \leq vk \text{ and } \exists x : vh[x] < vk[x]$
- $vh \parallel vk \Leftrightarrow \neg(vh < vk) \wedge \neg(vk < vh)$

# Vector Time – Basic Properties

---

- Isomorphism
- Strong Consistency
- Event Counting

# Vector Time – Basic Properties

## Isomorphism

- if events in a distributed system are timestamped using a system of vector clocks, following property holds:
  - if 2 events  $x$  and  $y$  have timestamps  $vh$  and  $vk$ , respectively, then
$$x \rightarrow y \Leftrightarrow vh < vk$$
$$x \parallel y \Leftrightarrow vh \parallel vk$$
- there is an isomorphism between the set of partially ordered events produced by a distributed computation and their vector timestamps
- very powerful, useful, and interesting property of vector clocks

# Vector Time – Basic Properties

## Isomorphism

- if the process at which an event occurred is known, test for comparing 2 timestamps can be simplified as:
  - if events  $x$  and  $y$  respectively occurred at processes  $p_i$  and  $p_j$  and have timestamps  $vh$  and  $vk$  respectively, then

$$x \rightarrow y \Leftrightarrow vh[i] \leq vk[i]$$

$$x \parallel y \Leftrightarrow vh[i] > vk[i] \wedge vh[j] < vk[j]$$

# Vector Time – Basic Properties

---

## Strong Consistency

- system of vector clocks is strongly consistent
- examination of vector timestamps of 2 events can determine if the events are causally related
- dimension of vector clocks can't be less than  $n$  for this property to hold

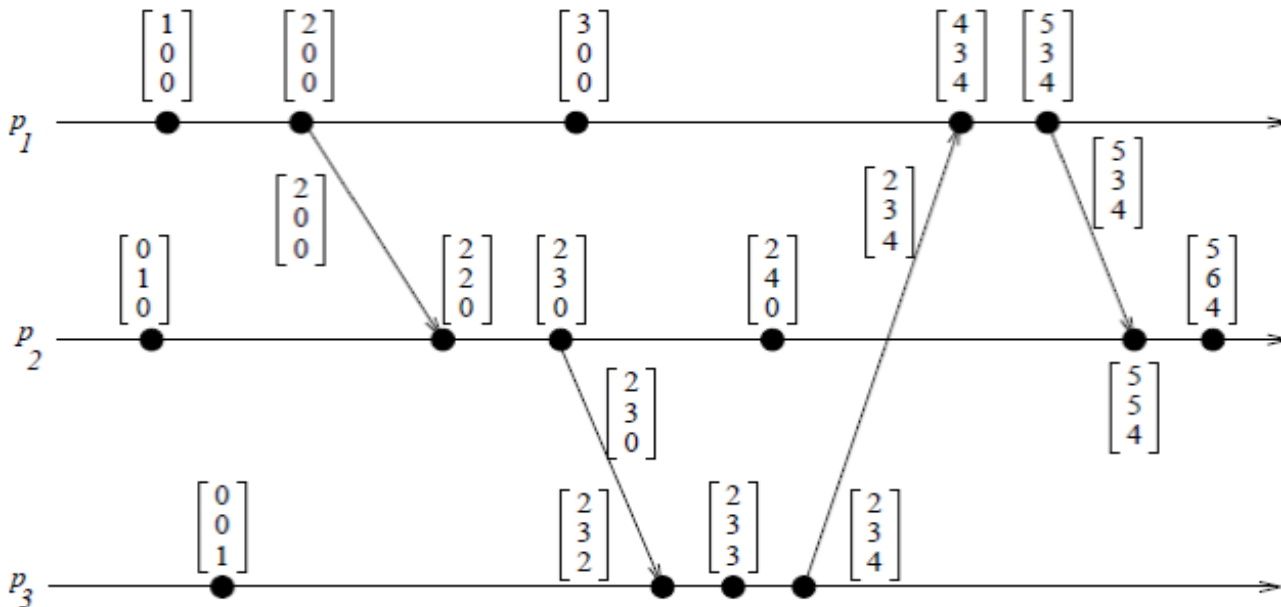
( $n$  = total no. of processes in the distributed computation)



# Vector Time – Basic Properties

## Event counting


- if  $d$  is always 1 in rule **R1** and  $vt_i[i]$  is the  $i^{th}$  component of vector clock at process  $p_i$ 
  - then  $vt_i[i] = \text{no. of events that have occurred at } p_i \text{ until that instant}$



# Vector Time – Basic Properties

## Event counting

- if an event  $e$  has timestamp  $vh$ 
  - $vh[j]$  = number of events executed by process  $p_j$  that causally precede  $e$
  - $\sum vh[j] - 1$  : total no. of events that causally precede  $e$  in the distributed computation



# Efficient Implementation of Vector Clocks



## Why is efficient implementation necessary?

- when no. of processes in a distributed computation is large
  - vector clocks will require piggybacking of huge amount of information in messages
  - messages are required for disseminating time progress and updating clocks
  - message overhead grows linearly with the no. of processes
  - message overhead is not affected by the no. of events occurring at the processors

# Singhal–Kshemkalyani's Differential Technique



- ❑ based on the observation - between successive message sends to the same process, only a few entries of the vector clock at the sender process are likely to change
- ❑ this is more likely when the number of processes is large
  - ❑ reason: only a few of them will interact frequently by passing messages
- ❑ fundamental idea behind the technique
  - ❑ when process  $p_i$  sends a message to process  $p_j$ , it piggybacks only those entries of its vector clock that differ since the last message sent to  $p_j$

# Singhal–Kshemkalyani's Differential Technique



- if entries  $i_1, i_2, \dots, i_{n_1}$  of the vector clock at  $p_i$  have changed to  $v_1, v_2, \dots, v_{n_1}$ , respectively, since the last message sent to  $p_j$ ,
  - then  $p_i$  piggybacks a compressed timestamp of the form
$$\{(i_1, v_1), (i_2, v_2), \dots, (i_{n_1}, v_{n_1})\}$$
to the next message to  $p_j$
- on receiving this message,  $p_j$  updates its vector clock as follows:
$$vt_j[i_k] = \max(vt_j[i_k], v_k) \text{ for } k = 1, 2, \dots, n_1$$

# Singhal–Kshemkalyani's Differential Technique



- **Benefit:**
  - reduces message size, communication bandwidth and buffer (to store messages) requirements
- **Worst case:**
  - every element of the vector clock has been updated at  $p_i$  since the last message sent to  $p_j$
  - next message from  $p_i$  to  $p_j$  will need to carry the entire vector timestamp of size  $n$
- **Average case: size of the timestamp on a message will be less than  $n$**

# Singhal–Kshemkalyani's Differential Technique




- **Requirements for implementation:**

- each process must remember the vector timestamp in the message last sent to every other process
- communication channels follow FIFO discipline for delivery of messages

# Singhal–Kshemkalyani's Differential Technique



- in general, the timestamp is of the form:  
 $\{(p_1, latest\_value), (p_2, latest\_value), \dots\}$   
 $p_i$  indicates  $p_i^{th}$  component of the vector clock has changed



- 2nd message from  $p_3$  to  $p_2$  contains a timestamp  $\{(3, 2)\}$
- informs  $p_2$  that the 3rd component of the vector clock has been modified and the new value is 2



# Singhal–Kshemkalyani's Differential Technique



## Benefits:

- cost of maintaining vector clocks in large systems can be substantially reduced
- especially if the process interactions exhibit temporal or spatial localities
- useful in applications
  - causal distributed shared memories
  - distributed deadlock detection
  - enforcement of mutual exclusion and localized communications

# Fowler–Zwaenepoel's Direct-Dependency Technique




- reduces the size of messages by transmitting only a scalar value
- no vector clocks are maintained on-the-fly
- process only maintains information regarding direct dependencies on other processes
- vector time for an event represents transitive dependencies on other processes
- vector time for event is constructed off-line from a recursive search of the direct dependency information at processes

# Fowler–Zwaenepoel's Direct-Dependency Technique



- $p_i$  maintains a dependency vector  $D_i$
- initially,  $D_i[j] = 0$  for  $j = 1, \dots, n$
- $D_i$  is updated using the following rules:
  1. When an event occurs at  $p_i$ ,  $D_i[i] = D_i[i] + 1$  (i.e., increment the vector component corresponding to its own local time by one)
  2. When  $p_i$  sends a message to  $p_j$ , it piggybacks the updated value of  $D_i[i]$  in the message
  3. When  $p_i$  receives a message from  $p_j$  with piggybacked value  $d$ ,  $p_i$  updates its dependency vector as follows:  $D_i[j] = \max\{D_i[j], d\}$

# Fowler–Zwaenepoel's Direct-Dependency Technique




# Fowler–Zwaenepoel's Direct-Dependency Technique



- $D_i$  reflects only direct dependencies
- at any instant,  $D_i[j]$  denotes the sequence no. of the latest event on  $p_j$  that **directly** affects the current state
- this event may precede the latest event at  $p_j$  that *causally* affects the current state

# Fowler–Zwaenepoel's Direct-Dependency Technique



- when  $p_4$  sends a message to  $p_3$ , it piggybacks a scalar that indicates the direct dependency of  $p_3$  on  $p_4$  because of this message
- then  $p_3$  sends a message to  $p_2$  piggybacking a scalar to indicate the direct dependency of  $p_2$  on  $p_3$  because of this message
- $p_2$  is indirectly dependent on  $p_4$  since  $p_3$  is dependent on  $p_4$

# Fowler–Zwaenepoel's Direct-Dependency Technique



- transitive (indirect) dependencies are not maintained by this method
- transitive dependencies
  - can be obtained only by recursively tracing the direct dependency vectors of the events off-line
  - involves computational overhead and latencies
- this method is ideal only for those applications
  - that do not require computation of transitive dependencies on the fly
  - eg. applications: causal breakpoints, asynchronous checkpoint recovery

# Fowler–Zwaenepoel's Direct-Dependency Technique

---



- saves cost considerably
- not suitable for applications that require on-the-fly computation of vector timestamps



# Physical Clock Synchronization



- no global clock or common memory
- each processor has its own internal clock and its own notion of time
- clocks can drift apart by several seconds per day, accumulating significant errors over time
- clock rates are different, may not remain always synchronized
- for most applications and algorithms that run in a distributed system, need to know time in one or more contexts:
  - time of the day at which an event happened on a specific machine in the network
  - time interval between two events that happened on different machines in network
  - relative ordering of events that happened on different machines in network

# Physical Clock Synchronization




- Clock synchronization –
  - ensuring that physically distributed processors have a common notion of time
  - performed to correct clock skew in distributed systems
  - synchronized to an accurate real-time standard like UTC (Universal Coordinated Time)
- Offset –
  - difference between the time reported by a clock and the real time
  - offset of the clock  $C_a$  is given by  $C_a(t) - t$  ( $C_a(t)$  – time of clock)
  - offset of clock  $C_a$  relative to  $C_b$  at time  $t \geq 0$  is given by  $C_a(t) - C_b(t)$

# Network Time Protocol



- uses offset delay estimation method
- involves a hierarchical tree of time servers
- primary server at the root synchronizes with the UTC
- next level contains secondary servers, which act as a backup to the primary server
- at the lowest level is the synchronization subnet which has the clients

# Network Time Protocol



- Each NTP message includes the latest three timestamps  $T_1$ ,  $T_2$ , and  $T_3$
- $T_4$  is determined upon arrival
- Peers A and B can independently calculate delay and offset using a bidirectional message stream

- $T_1, T_2, T_3, T_4$  - values of the 4 most recent timestamps as shown
- assume that clocks A and B are stable and running at the same speed
- $a = T_1 - T_3$ ,  $b = T_2 - T_4$
- if the network delay difference from A to B and from B to A, called differential delay, is small, the clock offset  $\theta$  and roundtrip delay  $\delta$  of B relative to A at time  $T_4$

$$\theta = (a+b)/2$$

$$\delta = a - b$$

# Network Time Protocol



- A pair of servers exchange pairs of timing messages.
- A store of data is then built up about the relationship between the two
- Servers (pairs of offset and delay).
- Each peer maintains pairs  $(O_i, D_i)$ , where:
  - $O_i$  – measure of offset ( $\theta$ )
  - $D_i$  – transmission delay of two messages ( $\delta$ )
- Offset corresponding to the minimum delay is chosen
- Assume that message  $m$  takes time  $t$  to transfer and  $m'$  takes  $t'$  to transfer
- Offset between  $A$ 's clock and  $B$ 's clock is  $O$
- If  $A$ 's local clock time is  $A(t)$  and  $B$ 's local clock time is  $B(t)$ 
  - $A(t) = B(t) + O$
- Then,  $T_{i-2} = T_{i-3} + t + O$ ,  $T_i = T_{i-1} - O + t$

# Network Time Protocol



- Assume  $t = t'$
- Offset  $O_i$  can be estimated as  $O_i = (T_{i-2} - T_{i-3} + T_{i-1} - T_i) / 2$
- Round-trip delay is estimated as  $D_i = (T_i - T_{i-3}) - (T_{i-1} - T_{i-2})$
- 8 most recent pairs of  $(O_i, D_i)$  are retained
- Value of  $O_i$  that corresponds to minimum  $D_i$  is chosen to estimate  $O$

# Recap Quiz



1. If an event  $e$  has a timestamp  $h$  in scalar clock, then the minimum logical duration required before producing event  $e$  is  
(a)  $h$  (b)  $h+1$  (c)  $h-1$  (d)  $1/h$
2. If the height of an event  $b = 5$ , then how many events would have preceded  $b$  on the longest causal path ending at  $b$  if we use a scalar clock?  
(a) 4 (b) 5 (c) 3 (d) 2
3. In the tie breaking mechanism used in total ordering in scalar clocks, a lower process identifier implies \_\_\_\_ priority  
(a) Equal (b) lower (c) higher (d) none of the above
4. Which of the following is not a mechanism to represent logical time in distributed computing systems?  
(a) Scalar clock (b) vector clock (c) matrix clock (d) list clock
5. Consider the event counting in vector clocks. if an event  $e$  has timestamp  $vh$ , then the number of events executed by process  $p_j$  that causally precede  $e$  will be  
(a)  $vh[j]$  (b)  $vh[j-1]$  (c)  $vh[j]$  (d)  $vh[0]$
6. Let the total number of processes  $n = 8$ , in a distributed environment. Then the dimension of the vector clocks can't be \_\_\_\_\_ 8 for strong consistency  
(a)  $<$  (b)  $>$  (c)  $\leq$  (d)  $\geq$
7. Which of the following would ensure the efficient implementation of vector clocks if the process interactions exhibit temporal or spatial localities in a distributed system?  
(a) Singhal-Kshemkalyani (b) Fowler-Zwaenepoel (c) physical clock synchronization (d) NTP
8. While implementing NTP in distributed computing systems, if  $A$ 's local clock time is  $A(t)$  and  $B$ 's local clock time is  $B(t)$ , then the Offset  $O$  is  
(a)  $A(t) + B(t)$  (b)  $A(t) * B(t)$  (c)  $A(t)/B(t)$  (d)  $A(t) - B(t)$
9. Fowler-Zwaenepoel approach does not maintain \_\_\_\_ dependencies  
(a) Commutative (b) associative (c) transitive (d) none of the above
10. In distributed computing systems, the clock skew can be corrected by \_\_\_\_\_ clock synchronization  
(a) Scalar (b) vector (c) matrix (d) physical

# Recap Quiz - key



Q1	Q2	Q3	Q4	Q5	Q6	Q7	Q8	Q9	Q10
c	b	c	d	a	a	a	d	c	d



# Major References



- Ajay D. Kshemkalyani, and Mukesh Singhal, Chapter 3, “Distributed Computing: Principles, Algorithms, and Systems”, Cambridge University Press, 2008 (Reprint 2013).